

std::execution Needs More Time

Document Number:	D4007R0
Date:	2026-02-02
Audience:	SG1, LEWG
Reply-to:	Vinnie Falco vinnie.falco@gmail.com Mungo Gill mungo.gill@me.com

Table of Contents

Revision History

Abstract

1. Introduction

2. The I/O Use Case

3. The Problem

 3.1 The Cost of Uncontrolled Allocation

 3.2 HALO Cannot Help

 3.3 `operator new` Is the Only Interception Point

 3.4 The `allocator_arg` Workaround

 3.5 The Receiver Arrives Too Late

4. Can It Be Fixed Later?

 4.1 The Gap Is Structural

 4.2 `await_transform` Cannot Help

 4.3 The Two-Tier Design Cannot Propagate Allocators

 4.4 P3826R3's Solutions Target a Different Problem

 4.5 ABI Lock-In Makes the Architectural Gap Permanent

5. The Committee's Own Record

 5.1 The Allocator Poll

 5.2 Post-Approval Changes

 5.3 The C++23 Precedent

6. A Path Forward

6.1 Avoiding API Lock-in

6.2 What Deferral Means

6.3 Who Bears the Cost

6.4 I/O Is in Scope

7. Conclusion

Appendix A - Code Examples

A.1 Why HALO Cannot Help

A.2 The `allocator_arg` Propagation Chain

A.3 The `connect / start / operator new` Sequence

A.4 Sender/Receiver vs Coroutine - Side by Side

Appendix B - Counterarguments Examined

B.1 P3826R3 and Algorithm Dispatch

B.2 Why `await_transform` Cannot Close the Gap

B.3 The Two-Tier Design and Allocator Propagation

Appendix C - Committee Record

C.1 The Single Async Model Debate

C.2 The C++23 Deferral

C.3 SG4 Networking Mandate

C.4 P3552 Design Approval and Forwarding

C.5 Hagenberg Design Reworks

C.6 P3796R1 and the Allocator Polls

C.7 Kona Algorithm Customisation Straw Poll

C.8 Outstanding LWG Defects

Appendix D - Post-Approval Modification Catalogue

Removals

Major Design Reworks

Wording Fixes and Corrections

Missing Functionality

Post-Adoption Issues

LWG Issues

Allocator-Related NB Comments (Kona 2025)

References

Revision History

- R0 - Initial revision.
-

Abstract

`std::execution` brings valuable ideas to C++ - structured concurrency, composable task graphs, and a clean separation of work description from execution policy. However, a fundamental workflow-ordering gap prevents coroutine-based I/O from using stateful allocators: `operator new` runs before the receiver exists, so the allocator arrives too late. This gap is architectural, not a missing feature that can be extended stably without destabilising the API. We recommend deferring `std::execution` to C++29, or alternatively publishing it as a Technical Specification, so that the allocator story can be resolved before the ABI is frozen.

1. Introduction

P2300R10 (<https://wg21.link/p2300r10>) (“`std::execution`”) represents years of careful work and genuine progress in structured asynchronous programming for C++:

- **Structured concurrency** with well-defined lifetime guarantees for asynchronous operations
- **Composable sender/receiver pipelines** that separate work description from execution policy
- **Schedulers and execution contexts** that give callers control over where work runs
- `run_loop` and foundational primitives that enable deterministic testing of async code
- **A formal model** for reasoning about asynchronous completion signatures

These are real achievements, and the community of users and implementers who have invested in sender/receiver architectures - NVIDIA CCCL, stdexec, libunifex, Folly - has demonstrated substantial value in CPU-bound parallelism and GPU dispatch.

The concern this paper raises is narrower than a general objection: `std::execution` works well for its primary use cases. The gap arises when a sender pipeline must select a custom allocation strategy. The allocator must be available before construction, but the sender/receiver model provides it only after the frame has already been created.

Couroutine-based I/O makes this gap acute. Asynchronous execution makes optimising away the coroutine frame allocation unlikely or impossible, and the impact of multiple extra allocations per I/O operation can be disastrous. The committee has explicitly placed this use case in scope (see Section 6.4); the standard must accommodate it.

The question is whether the gap can be closed without either (1) a breaking API change or (2) leaving the current interface as vestigial cruft. Section 4 examines each proposed escape route; unfortunately, the authors have been unable to find a way to close this gap.

The allocator timing gap is not a missing feature that can be added later; it is architectural: `operator new` runs before `connect()`, and no library-level change can alter that sequencing without destabilising the API. Discovering this gap before standardisation is fortunate, not a failure. The feature's size and complexity make thorough review genuinely difficult, and it is not surprising that a gap visible only in the I/O use case was not identified earlier. There is still time to address it.

2. The I/O Use Case

`std::execution` aspires to be a general-purpose asynchronous framework, including for networked I/O. Here is what that use case requires.

A typical networking application launches coroutines that perform I/O:

```

// User's coroutine - handles one client connection
task<void> handle_client(tcp_socket socket)
{
    char buf[1024];
    auto [ec, n] = co_await socket.async_read(buf);
    if (n > 0)
        process(buf, n); // data may accompany eof
}

// Application launches the coroutine
int main()
{
    io_context ctx;
    tcp_socket sock = accept_connection(ctx);
    auto alloc = ctx.get_frame_allocator();

    run_async(ctx.get_executor(), alloc)(
        handle_client(std::move(sock)) // Frame allocated HERE
    );

    ctx.run();
}

```

The critical timing:

```

run_async(executor, alloc)(handle_client(sock))
|
    +-> promise_type::operator new runs HERE
        Allocator must already be known

```

When `handle_client(sock)` is evaluated, the compiler calls `promise_type::operator new` to allocate the coroutine frame. The allocator must be available at this moment - before the coroutine body executes, before any `co_await`, before any connection to a receiver.

3. The Problem

The allocator timing gap has three facets: the cost of uncontrolled allocation, the absence of any interception mechanism besides `operator new`, and the ordering constraint that places allocation before the receiver exists.

3.1 The Cost of Uncontrolled Allocation

High-performance I/O servers cannot tolerate uncontrolled heap allocation. Realistic server code calls nested coroutines for parsing, processing, and serialisation - each `task<T>` call allocates a frame:

```
task<void> handle_connection(tcp_socket socket)
{
    char buf[1024];
    while (socket.is_open()) {
        auto n = co_await socket.async_read(buf); // I/O awaitable: no frame
        co_await process_request(socket, buf, n); // task<void>: frame allocated
    }
}
// handle_connection itself: 1 frame per connection
// process_request (and coroutines it calls): frame(s) per request
// 10,000 connections x 1,000 requests/sec = 10 million+ frame allocations/sec
// global heap contention becomes the bottleneck
```

The solution is **stateful allocators**:

- recycling allocators that cache recently-freed frames
- arena allocators tied to connection scope
- pool allocators optimised for common frame sizes
- `pmr::memory_resource` implementations that customise the upstream allocation strategy

Modern general-purpose allocators (jemalloc, mimalloc, tcmalloc) achieve impressive throughput via per-thread caching, but they solve a different problem. They reduce contention on the global heap; they do not address the allocation patterns that I/O servers produce. Coroutine frames are short-lived, uniformly sized, and frequently allocated on one thread but freed on another - a cross-thread pattern that defeats per-thread caches.

More fundamentally, a general-purpose allocator has no notion of application-level context. Different coroutine chains might require different policies. A multi-tenant server may enforce per-tenant memory limits, with each tenant's connections using an allocator bound to that tenant's quota. A connection-scoped arena can reclaim all frames in one operation when the connection closes. These policies require the stateful allocators listed above - and that state must be accessible when `operator new` executes.

3.2 HALO Cannot Help

HALO (Heap Allocation eLision Optimisation) allows compilers to elide coroutine frame allocation when the frame's lifetime is provably bounded by its caller. For I/O coroutines launched onto an execution context, the frame outlives its caller - the caller returns immediately after launching the operation, and the frame persists until the OS signals completion. The compiler cannot prove bounded lifetime when the lifetime depends on an external event; hence, HALO cannot apply in this common case, and allocation is mandatory. See Appendix A.1 for a deeper analysis and code example.

It is also worth noting that HALO does nothing when the coroutine's implementation is not visible.

3.3 `operator new` Is the Only Interception Point

The senders/receivers facility, as defined today, provides exactly one mechanism to intercept coroutine frame allocation:

```
struct promise_type
{
    static void* operator new(std::size_t size)
    {
        // This is the ONLY place to intercept allocation
        // The allocator must be known RIGHT NOW

        std::pmr::memory_resource* mr = /* ??? */;
        return mr->allocate(size, alignof(std::max_align_t));
    }
};
```

The allocator - along with its state - must be discoverable at this exact moment. No mechanism that provides the allocator later can help. The frame is already allocated by the time any later mechanism executes.

3.4 The `allocator_arg` Workaround

P3552R3 (<https://wg21.link/p3552r3>) (“Add a Coroutine Task Type”) provides a mechanism for passing an allocator via a parameter to a given coroutine via `std::allocator_arg_t`. This additional syntax solves the initial allocation: the caller passes the allocator explicitly at the call site. However, propagation through call chains is unsolved. Every coroutine in the chain must accept variadic template arguments, query the allocator from its environment, and explicitly forward that allocator to every child coroutine call. The allocator becomes viral - function signatures - now with two additional parameters - no longer reflect algorithmic intent, and forgetting to forward the allocator silently breaks the chain. See Appendix A.2 for the complete propagation example.

Allocator propagation is admittedly a general challenge in C++, not unique to sender/receiver. However, the standard library addresses it for containers through `uses_allocator` construction and `scoped_allocator_adaptor`, which give generic code a uniform propagation path. No equivalent exists for coroutine frames, and the ordering constraint described in Section 3.5 - the receiver arrives after `operator new` - means that the natural source of the allocator in a sender/receiver pipeline is structurally unavailable at the point of allocation.

3.5 The Receiver Arrives Too Late

Context in the sender/receiver model in `std::execution` flows backward from receiver to sender. This creates an ordering gap for coroutine frame allocation:

```
task<int> my_coro(); // Coroutine returning a sender

auto sndr = my_coro(); // Step 1: operator new runs HERE
// Frame already allocated

auto op = connect(std::move(sndr), rcvr); // Step 2: Receiver connected HERE
// get_allocator(get_env(rcvr))
// available NOW - too late

start(op); // Step 3: Operation starts
```

P2300R4 (<https://wg21.link/p2300r4>) (“`std::execution`,” 2022-01-18) acknowledges this timing:

“In the sender/receiver model... contextual information is associated with the receiver and is queried by the sender and/or operation state **after** the sender and the receiver are `connect`-ed.”

P3826R3 (<https://wg21.link/p3826r3>) (2026-01-05) confirms the consequence:

“The receiver is not known during early customization. Therefore, early customization is irreparably broken.”

The allocator is part of this contextual information. It becomes available only after `connect()` - but the coroutine frame was allocated in Step 1, before `connect()` was called.

For coroutine-based I/O, the allocator arrives too late. This lack of customisability is not a flaw in the sender/receiver formalism itself; rather, it is an unanticipated gap that emerges only when that formalism meets the concrete requirements of coroutine-frame allocation.

4. Can It Be Fixed Later?

How we proceed will depend on our ability to address the issue in a way that does not require either (1) an API/ABI break or (2) leaving what we have as a vestigial remnant of an incomplete framework.

Each potential escape route of which the authors are currently aware - `await_transform`, the two-tier design, P3826R3’s algorithm dispatch fixes - leaves the fundamental sequencing constraint intact.

4.1 The Gap Is Structural

`operator new` runs before `connect()`. No library-level change can alter this sequencing - it is determined by the C++ language specification for coroutines’ interacting with the sender/receiver protocol. There is no post-hoc extension point between “coroutine frame is allocated” and “receiver is connected.” The only way to close the gap is to change the sequencing - which means changing the API.

4.2 `await_transform` Cannot Help

A natural question is whether `promise_type::await_transform` could inject the allocator into child coroutine calls, eliminating explicit `allocator_arg` passing. It cannot. In `co_await child_coro(args...)`, the function call `child_coro(args...)` is evaluated first - the child’s `operator new` fires and the frame is allocated before `co_await` processing begins. By the time `await_transform` sees the returned `task<T>`, the child’s frame has already been allocated without the parent’s allocator. See Appendix B.2 for a deeper evaluation-order analysis.

4.3 The Two-Tier Design Cannot Propagate Allocators

P3552R3 (<https://wg21.link/p3552r3>) establishes a two-tier model: the allocator is a creation-time concern passed at the call site, while the scheduler and stop token are connection-time concerns from the receiver. The promise's `get_env()` exposes the stored allocator to child operations via `get_allocator`. That mechanism gives callers a standard way to allocate the initial coroutine frame with a chosen allocator - a valuable contribution to the design.

Propagating that choice through a chain of nested coroutine calls, however, remains the harder and (as yet) unsolved problem.

With standard containers, the allocator type is part of the type signature (`vector<T, Alloc>`), and `uses_allocator` construction gives generic code a standard way to propagate the allocator through nested construction. With coroutines, the allocator is consumed inside `operator new` (which runs before the receiver exists) and the coroutine's return type (`task<T>`) does not carry the allocator type. There is no `uses_allocator` equivalent for coroutine frame allocation. Nor is it clear that one could be created: any indirection layer that wraps a coroutine invocation risks defeating Heap Allocation eLision Optimisation (HALO), because the compiler must see the coroutine call directly at the call site.

The propagation burden is compounded by the calling convention: every call site must pass both `std::allocator_arg` and the allocator object as leading arguments. A coroutine that naturally takes a single argument becomes a three-parameter call, two-thirds of which is allocator plumbing rather than application logic - creating strong pressure to simply omit the allocator and fall back to the global heap.

4.4 P3826R3's Solutions Target a Different Problem

P3826R3 (<https://wg21.link/p3826r3>) ("Fix Sender Algorithm Customization," 2026-01-05) proposes important fixes for sender algorithm customisation - a real problem prompted by NB comment US 207 (<https://github.com/cplusplus/nbballot/issues/903>) and resolved over three revisions of the proposal. However, its five proposed solutions all target algorithm dispatch - i.e., which implementation of `then`, `let_value`, or `bulk` should run. Four of the five do not change when the allocator becomes available. The fifth - remove all `std::execution` - resolves the gap by deferral. See Appendix B.1 for a deeper analysis of each solution.

4.5 ABI Lock-In Makes the Architectural Gap Permanent

Once standardised, the relationship between `operator new` and `connect()` becomes part of the ABI. A fix would likely need `connect()` to propagate allocator context before the coroutine frame is allocated - a structural change to the sender protocol. If such a change proves ABI-incompatible, the alternative is a second async framework alongside the first. Neither outcome is impossible to manage, but both carry long-term costs that are worth weighing before the ABI boundary is set.

5. The Committee's Own Record

The committee's own proceedings confirm that the allocator story is unresolved and the design is still actively evolving.

5.1 The Allocator Poll

P3796R1 (<https://wg21.link/p3796r1>) (Dietmar Kuhl, “Coroutine Task Issues”) was reviewed by LEWG in a September 2025 telecon. The allocator question was polled directly:

“We would like to use the allocator provided by the receivers env instead of the one from the coroutine frame”

SF	F	N	A	SA
0	0	5	0	0

Attendance: 14. Outcome: strictly neutral.

The entire room abstained to neutral - an outcome consistent with the structural difficulty described in Section 4.3: without a mechanism to propagate allocator context through nested coroutine calls, the committee had no direction to endorse. The author was directed to explore preferred solutions with code examples.

Two US national body comments confirm the same concern. US 255 (LWG4335) requests that `task` use the allocator from the receiver's environment rather than the one supplied at coroutine creation. The NB comment itself acknowledges the difficulty: “the used allocator type needs to be known when the coroutine is created. At that time the receiver isn't known, yet.” The core issue is that coroutine `operator new` overload resolution requires the allocator type at compile time - a runtime pointer to a `pmr::memory_resource` is not sufficient. At Kona (2025-11-06) LEWG polled the direction:

"We approve the direction proposed in the NB comment: US 255-384 [task.promise] Use allocator from receiver's environment LWG4335 and ask for wording."

SF	F	N	A	SA
4	7	0	1	0

Attendance: 17. Outcome: consensus in favour.

The sole A vote - from the task paper's author - noted that using the receiver's allocator "goes against the general behaviour of allocators." Wording remains outstanding. US 253 (LWG4333) requests arbitrary allocator support for the coroutine frame; wording is also still required.

LWG 4356 (`connect()` should use `get_allocator(get_env(rcvr))`) confirms the gap has been filed as a specification defect - not merely an external complaint, but a committee-level acknowledgement of the timing problem.

5.2 Post-Approval Changes

The volume of papers written to fix, rework, or complete `std::execution` after its approval for C++26 is extraordinary. Mature designs do not require this level of post-approval remediation.

Period	Removals	Reworks	Wording	Missing Functionality	LWG	Total
Pre-Wroclaw (Mar-Oct 2024)	1	5	-	1	-	7
Pre-Hagenberg (Nov 2024-Feb 25)	1	-	2	2	3	8
Pre-Sofia (Mar-Jun 2025)	-	2	-	7	1	10
Pre-Kona (Jul-Nov 2025)	-	3	3	1	7	14
Pre-London (Dec 2025-Feb 2026)	-	1	1	-	-	2
Total	2	11	6	11	11	41

See Appendix D for the complete listing with dates, authors, and current status.

Key observations:

- **The churn is accelerating, not slowing.** The pre-Kona period produced 14 items - the highest volume - including Priority 1 safety defects. Two more papers appeared in the January 2026 mailing, with the London meeting still months away.
- **The severity has not decreased.** The pre-Kona period includes Priority 1 LWG defects: a dangling-reference vulnerability in `transform_sender` (LWG 4368) and an unconstrained alias in `connect_result_t` (LWG 4206). Two outstanding NB comments on allocator support remain without wording.
- **Design reworks span the entire timeline since `std::execution` was voted into the C++26 working draft (Tokyo, March 2024).** From P2855R1 (replacing `tag_invoke`, March 2024) through P3927R0 (fixing `task_scheduler` bulk execution, January 2026) - 11 rework papers over 22 months.

In total: **30 papers, 11 LWG defects, and 2 NB comments - 43 items modifying a single feature after its approval.**

5.3 The C++23 Precedent

This is not the first time maturity concerns have led to deferral. P2300 failed to achieve consensus for C++23:

LEWG 2022-01-12: “Produce a P2300R4, modified as described below, and then send the revised paper to LWG for C++23 with priority 1”

SF:15 / F:9 / N:1 / A:3 / SA:6 (Attendance: 39)

Recorded SA reasons included “maturity/readiness” and “Feels like ranges all over again. Needs maturity.” The same pattern of ongoing design changes - algorithm reworks, allocator timing unresolved, multiple LWG issues opened during review - is present again. The difference is that the concerns which were speculative in 2022 are now concrete: specific LWG issues, neutral committee polls, and unresolved NB comments document the gaps rather than merely anticipate them.

6. A Path Forward

This paper asks the committee to consider whether the allocator timing gap should be resolved before the `std::execution` API is frozen in the IS.

6.1 Avoiding API Lock-in

The P2300 authors are best placed to propose a resolution to the allocator timing gap. We welcome any solution that closes it - whether by changes to the `task` type, to the coroutine promise protocol, or to the interaction between `operator new` and `connect`. Our concern is that freezing the API in the IS before the gap is addressed would lock C++ into an interface that cannot serve coroutine-based I/O without work-arounds.

If no solution is forthcoming in the C++26 time frame, deferral to a later revision of the standard would avoid ABI lock-in while preserving the committee's freedom to fix what remains.

6.2 What Deferral Means

Deferring `std::execution` from C++26 does not discard the work done by the P2300 authors and implementers. Production-quality implementations already exist and are in active use. Deferral means the committee does not freeze the ABI until the remaining architectural gaps - allocator timing, allocator propagation, and the interaction with coroutine task types - are resolved.

Deferral is not failure. The committee deferred trivial relocation to C++29 because open design questions remained; the result was a stronger proposal built on the foundations already laid. The same precedent applies here: the P2300 authors' investment is preserved in full, and production-quality implementations (`stdexec`, NVIDIA CCCL, `libunifex`, Folly) remain available to users regardless of the standardisation timeline. A Technical Specification is a further option, though existing implementations already serve that role in practice.

6.3 Who Bears the Cost

The urgency question deserves an honest answer: who is harmed by deferral, and who is harmed by premature standardisation?

Organisations building on `std::execution` today are not blocked by the absence of an IS specification - they build against concrete library implementations. Unlike language features, non-standard library implementations are straightforward to adopt; organisations can use Boost, `stdexec`, or other implementations without compiler extensions. If the committee were to defer this feature until a future iteration of the language - whether or not it is released as a TS - no current users would be affected.

The networking use case - which arguably represents the largest constituency for asynchronous C++ - is the community most affected by shipping `std::execution` that has an incomplete allocator story. SG4 mandated that networking be built exclusively on sender/receiver (see Section 6.4). If the current sender/receiver framework ships without a workable allocator propagation model for coroutine-based I/O, the networking community faces a choice between inefficient workarounds and non-standard alternatives.

This framing is not an accusation that the committee is neglecting networking; rather, it is an observation that the current timeline favours constituencies that do not need allocator propagation (GPU dispatch, CPU-bound parallelism) over the constituency that does (I/O-bound networking). Deferral preserves optionality for everyone; premature standardisation imposes a concrete cost on the use case with the broadest demonstrated demand.

6.4 I/O Is in Scope

A likely response is that `std::execution` is for structured concurrency and parallelism, and networking is a separate paper track. The committee's own record contradicts this assertion.

SG4 polled at Kona (November 2023) on [P2762R2](https://wg21.link/p2762r2) (“Sender/Receiver Interface For Networking”):

“Networking should support only a sender/receiver model for asynchronous operations; the Networking TS’s executor model should be removed”

SF:5 / F:5 / N:1 / A:0 / SA:1 - Consensus.

Networking must be built on `std::execution`’s sender/receiver model. The task type ([P3552R3](https://wg21.link/p3552r3) (`co_await socket.async_read(buf)`)) is the primary I/O usability layer - it exists so that users can write If the coroutine integration cannot access the allocator, the I/O use case is effectively unserved. See Appendix C.1-C.3 for the complete committee record on this question.

7. Conclusion

`std::execution` represents important progress in structured asynchronous programming for C++. Its contributions - structured concurrency, composable pipelines, and a formal model for async completion - are valuable and worth building upon.

For coroutine-based asynchronous I/O, however, a fundamental timing gap remains:

1. **Allocation is mandatory** - HALO cannot optimise away frames that outlive their callers
2. **Stateful allocators are required** - 10 million+ allocations/sec demands recycling, pooling, or arena strategies
3. **`operator new` is the only interception point** - the allocator must be known when the frame is allocated

4. `std::execution` provides the allocator too late - the receiver's environment is available only after `connect()`
5. The gap cannot be fixed later - `operator new` runs before `connect()`, and no library change can alter that sequencing without destabilising the API

P2300R4 (<https://wg21.link/p2300r4>) established the sender/receiver context model in January 2022. Four years later, the design continues to evolve: P3826R3 (<https://wg21.link/p3826r3>) proposes significant architectural changes to `transform_sender`, removes early customization entirely, and restructures the relationship between `continues_on` and `schedule_from`. This level of active change is healthy for a maturing design, but it signals that the work is not yet finished.

If `std::execution` ships with C++26 but cannot serve coroutine-based I/O without workarounds, the committee risks standardising a framework that works for CPU-bound parallelism and GPU dispatch but leaves the largest async constituency - networking - reliant on non-standard alternatives.

The goal is not to block `std::execution`. It is to ensure that when it ships, it will be capable of serving the full spectrum of its intended use cases - including coroutine-based I/O - so that the C++ community gets one async framework, not two.

We ask the committee to address the allocator timing gap before freezing the `std::execution` API in the IS. We welcome a solution from the P2300 authors and are eager to collaborate on one. If the gap cannot be closed in the C++26 time frame, we encourage the committee to continue this work with a view to inclusion in C++29.

Appendix A - Code Examples

A.1 Why HALO Cannot Help

HALO (Heap Allocation eLision Optimisation) allows compilers to elide coroutine frame allocation when the frame's lifetime is provably bounded by its caller. When an I/O coroutine is launched onto an execution context, the frame must outlive the launching function:

```

task<size_t> read_data(socket& s, buffer& buf)
{
    co_return co_await s.async_read(buf); // Suspends for network I/O
}

void start_read(executor ex)
{
    start_detached(ex, read_data(sock, buf)); // Frame must outlive this function
                                                // Caller returns immediately
}

```

`start_detached` transfers the frame's lifetime to the execution context. The caller returns, and the frame persists until the OS signals I/O completion - which may take microseconds or seconds. The compiler cannot prove bounded lifetime, so HALO cannot apply and allocation is mandatory.

A note of precision: HALO cannot apply when the coroutine frame outlives its caller, which is the common case for I/O coroutines launched onto an execution context. In the nested case - where coroutine A `co_await`s coroutine B - A's frame remains alive while B executes, and B's frame lifetime is bounded by A's. In principle, a sufficiently advanced compiler could elide B's allocation into A's frame. No current compiler does this for `task<T>` coroutines, and the optimisation becomes infeasible once any frame in the chain is transferred to an execution context. HALO cannot help when frames outlive their callers - which is the case that matters for I/O.

A.2 The `allocator_arg` Propagation Chain

P3552R3 (<https://wg21.link/p3552r3>) provides a mechanism for passing an allocator to a single coroutine via `std::allocator_arg_t`. Here is what propagation through a call chain looks like in practice:

```

namespace ex = std::execution;

template<class... Args>
task<void> level_three(Args&&...) { co_return; }

template<class... Args>
task<void> level_two(int x, Args&&...) {
    auto alloc = co_await ex::read_env(ex::get_allocator); // Query
    co_await level_three(std::allocator_arg, alloc); // Forward
}

template<class... Args>
task<int> level_one(int v, Args&&...) {
    auto alloc = co_await ex::read_env(ex::get_allocator); // Query
    co_await level_two(42, std::allocator_arg, alloc); // Forward
    co_return v;
}

```

Every coroutine in the chain must accept the allocator argument (whether via variadic template parameters, defaulted arguments, or an explicit `allocator_arg` overload), query the current allocator, and forward it to each child. Function signatures no longer reflect algorithmic intent. Forgetting to forward the allocator silently breaks the chain. The allocator becomes viral - polluting interfaces throughout the code-base.

A.3 The `connect` / `start` / `operator new` Sequence

The complete timing sequence showing the gap between frame allocation and receiver connection:

```

auto sndr = my_coro(); // Step 1: operator new runs HERE
                      // Frame already allocated

auto op = connect(std::move(sndr), rcvr); // Step 2: Receiver connected HERE
                                         // get_allocator(get_env(rcvr))
                                         // available NOW - too late

start(op); // Step 3: Operation starts

```

P2300R4 (<https://wg21.link/p2300r4>) (2022-01-18) acknowledges this timing:

“In the sender/receiver model... contextual information is associated with the receiver and is queried by the sender and/or operation state **after** the sender and the receiver are `connect`-ed.”

P3826R3 (<https://wg21.link/p3826r3>) (2026-01-05) confirms the consequence for customisation:

“The receiver is not known during early customization. Therefore, early customization is irreparably broken.”

Note: the P3826R3 quote refers specifically to algorithm dispatch customisation, not to allocator timing. But the underlying cause is the same: the receiver - and its environment including the allocator - is not available until `connect()`.

A.4 Sender/Receiver vs Coroutine - Side by Side

The same `handle_connection` logic in both models illustrates where the gap opens.

Pure sender/receiver (no coroutines):

```
auto handle_connection(tcp_socket socket)
{
    return just(std::move(socket))
        | let_value([](tcp_socket& s) {
            return async_read(s)
                | let_value([&](auto data) {
                    return process(s, data);
                });
        });
}

// Usage:
auto op = connect(handle_connection(sock), rcvr);
//                                     ^^^^
// The receiver's allocator is available here.
// connect() can allocate the entire operation state
// tree using get_allocator(get_env(rcvr)).
start(op);
```

The standard sender algorithms (`just`, `let_value`, etc.) build a lazy description; no operation state is allocated until `connect(sender, receiver)`, at which point the receiver's environment - including the allocator - is available for the entire operation state tree. A custom sender theoretically could allocate earlier, but the protocol's design defers allocation to `connect()` precisely so that the receiver's allocator can be used. There is no timing gap.

Coroutine model:

```
task<void> handle_connection(tcp_socket socket)
{
    // operator new has ALREADY run by this point.
    // The frame is allocated. The receiver does not exist yet.
    char buf[1024];
    auto n = co_await socket.async_read(buf);
    co_await process_request(socket, buf, n);
}
```

The coroutine frame is allocated during the function call - before any receiver exists. This is the precise moment the gap opens. Any sender that allocates storage at creation time faces the same timing mismatch, but coroutines make it pervasive: `task<T>` always allocates a frame in `operator new` before any receiver exists. Since `task<T>` is the primary usability layer for writing sender/receiver code, the gap affects the most common authoring path.

Appendix B - Counterarguments Examined

B.1 P3826R3 and Algorithm Dispatch

P3826R3 (<https://wg21.link/p3826r3>) proposes important fixes for sender algorithm customisation, prompted by NB comment US 207 (<https://github.com/cplusplus/nbballot/issues/903>). The resolution went through three revisions before acceptance in January 2026 - evidence of ongoing design evolution, but not itself an indictment.

P3826 targets algorithm dispatch and was never intended to address allocator timing. Because its Solution 4.5 restructures `transform_sender` to accept the receiver's environment, however, a natural question is whether that change also closes the allocator gap. It does not - but examining why is instructive.

P3826 offers five solutions. All target algorithm dispatch - i.e., which implementation of `then`, `let_value`, or `bulk` should run:

Solution 4.1: Remove all `std::execution`. This would resolve the timing gap by deferral.

Solution 4.2: Remove customisable sender algorithms. This removes `then`, `let_value`, `bulk`, etc. It does not change when the allocator becomes available.

Solution 4.3: Remove sender algorithm customisation. This removes the ability to customise algorithms. It does not change when the allocator becomes available.

Solution 4.4: Ship as-is, fix via DR. This defers the fix. It does not change when the allocator becomes available.

Solution 4.5: Fix algorithm customisation now. P3826's recommended fix restructures `transform_sender` to take the receiver's environment, changing information flow at `connect()` time:

```
get_completion_domain<set_value_t>(get_env(sndr), get_env(rcvr))
```

This tells senders where they will start, enabling correct algorithm dispatch. It does not change when the allocator becomes available - the receiver's environment is still only queryable after `connect()`.

There was an assertion made during the Kona discussion that removing sender algorithm customisation could preclude shipping the task type entirely, since the task type's utility depends on the ability to customise algorithms. This would imply that the allocator timing issue, the algorithm customisation issue, and the task type are architecturally intertwined.

B.2 Why `await_transform` Cannot Close the Gap

A reader familiar with C++ coroutine machinery may ask whether `promise_type::await_transform` could inject the allocator into child coroutine calls, eliminating explicit `allocator_arg` passing.

The issue is evaluation order. In the expression:

```
co_await child_coro(args...)
```

The function call `child_coro(args...)` is evaluated first. The child's `operator new` fires, the coroutine frame is allocated, the promise is constructed, and `get_return_object()` runs - all during the function call expression. Only then does `co_await` processing begin, at which point `await_transform` sees the returned `task<T>`.

By the time `await_transform` executes, the child's frame has already been allocated without the parent's allocator. There is no customisation point in the C++ coroutine specification that allows a caller to inject context into a callee's allocation. `operator new`, `get_return_object`, and `initial_suspend` all execute during the function call expression, before `co_await` processing begins.

`await_transform` can inject context into a child's execution phase - this is how P3552's `task` propagates the scheduler and stop token via the receiver environment at `connect()` / `start()` time. But allocation and execution are different phases: allocation happens at call time, execution happens at start time. The sender/receiver model's backward-flow context is designed for the execution phase. It has no mechanism to influence the allocation phase, because allocation precedes the existence of the receiver.

B.3 The Two-Tier Design and Allocator Propagation

P3552R3 (<https://wg21.link/p3552r3>) establishes a two-tier model:

- **Creation-time concerns:** The allocator is passed at the call site via `std::allocator_arg_t`. The promise stores it and exposes it through `get_env()`.
- **Connection-time concerns:** The scheduler and stop token come from the receiver's environment, available after `connect()`.

This is a thoughtful design that correctly separates two categories of concern. It solves allocating the initial coroutine frame with a custom allocator - a real contribution.

The unsolved problem is propagation. When `level_one` calls `level_two` which calls `level_three`, each child coroutine needs the allocator at its call site - before it has any connection to a receiver. The parent must query its own environment, extract the allocator, and pass it explicitly.

A likely counterargument is that this is analogous to how allocators work for standard containers: you pass the allocator at construction time, and the container propagates it internally. But coroutines are fundamentally different:

- With containers, the allocator type is part of the type signature (`vector<T, Alloc>`). With coroutines, the return type is `task<T>` - it does not carry the allocator type.
- With containers, `uses_allocator` construction gives generic code a standard way to propagate the allocator. With coroutines, there is no equivalent mechanism.
- With containers, the allocator is passed once at construction and the container handles propagation internally. With coroutines, every call site in the chain must manually query and forward.

Until a propagation mechanism comparable to `uses_allocator` exists for coroutine frame allocation, the two-tier design solves only the first call in the chain. Production I/O code involves deep call chains, and the manual forwarding burden makes it impractical to maintain allocator discipline throughout.

One possible direction would be to make the promise type a factory for child coroutine calls, allowing it to inject the allocator automatically. This would require either a new coroutine customisation point that fires before the child's `operator new`, or a change to the evaluation order specified in [expr.await] so that frame allocation can be deferred until the parent's promise has had a chance to supply the allocator. Either path is a language-level change to the coroutine specification, and would need time to design and validate.

Appendix C - Committee Record

C.1 The Single Async Model Debate

On 2021-09-28 (attendance: 49), the committee debated whether `std::execution` should be the single async model for C++:

“We must have a single async model for the C++ Standard Library”

SF	F	N	A	SA
5	9	10	11	4

No consensus.

The committee could not agree on a single model, but the question itself reveals the ambition: `std::execution` was being positioned to serve all async use cases including I/O.

C.2 The C++23 Deferral

LEWG 2022-01-12:

“Produce a P2300R4 (`std::execution`), modified as described below, and then send the revised paper to LWG for C++23 with priority 1”

SF	F	N	A	SA
15	9	1	3	6

Attendance: 39. Weak consensus.

Recorded SA reasons included: “maturity/readiness/tag_invoke and the numerous change in the last few meetings,” “Feels like ranges all over again. Needs maturity,” and “Too few reviews from embedded people.”

The C++23 deferral happened because of maturity concerns. The same pattern of ongoing design changes is present again.

C.3 SG4 Networking Mandate

SG4 polled at Kona (November 2023) on P2762R2 (<https://wg21.link/p2762r2>):

“Networking should support only a sender/receiver model for asynchronous operations; the Networking TS’s executor model should be removed”

SF	F	N	A	SA
5	5	1	0	1

Consensus.

SG4’s feature requirements for networking (Kona 2023):

Feature	Yes	No
async_read	8	2
async_write	7	2
timeout	8	2
connect(context, name)	8	1

These are the I/O primitives that will be built on `std::execution`.

C.4 P3552 Design Approval and Forwarding

The P3552R1 design approval poll (LEWG, 2025-04-22):

SF	F	N	A	SA
5	6	6	1	0

Attendance: 23.

Six neutral votes matching the number of favourable votes is a notably soft result for a design approval poll.

The forwarding poll (LEWG, 2025-05-06):

“Forward P3552R1 with the action items discussed to LWG for C++29”

SF:5 / F:7 / N:0 / A:0 / SA:0 - unanimous.

“Forward P3552R1 to LWG with the action items discussed with a recommendation to apply for C++26 (if possible).”

SF:5 / F:3 / N:4 / A:1 / SA:0 - weak consensus, with “if possible” qualifier.

C++29 was unanimous; C++26 was conditional and weak.

C.5 Hagenberg Design Reworks

At Hagenberg (February 2025), SG1 forwarded P3552R0 with required architectural changes: replacing `continues_on` with a new `affine_on` mechanism, adding a cancellation channel, supporting errors without exceptions, and implementing scheduler affinity (SF:7 / F:0 / N:0 / A:1 / SA:0). The sole A voter specifically wanted “a more general concept for scheduler affinity to come to SG1 before we ship something in the standard.”

LEWG then polled on `affine_on` replacing `continue_on` (SF:9 / F:6 / N:2 / A:1 / SA:0) and on error-without-exceptions support (SF:6 / F:6 / N:4 / A:1 / SA:1). These are fundamental design changes - not editorial fixes - made in the same year as the C++26 deadline.

C.6 P3796R1 and the Allocator Polls

See Section 5.1 for the LEWG allocator poll. The broader context:

P3796R1 (<https://wg21.link/p3796r1>) (Dietmar Kuhl, “Coroutine Task Issues”) was reviewed in LEWG telecons during August-September 2025. Some sections achieved consensus; others remain pending. Multiple LWG issues were opened (4329-4332, 4344).

US 255 ([cplusplus/nbballot#959](https://github.com/cplusplus/nbballot/issues/959) (<https://github.com/cplusplus/nbballot/issues/959>), LWG4335): “Use allocator from receiver’s environment.” The NB comment states: “Normally, the `get_allocator` query forwards the allocator from the receiver’s environment. For task the `get_allocator` query used for `co_awaited` senders uses the allocator passed when creating the coroutine or the default if there was none. It should use the receiver’s environment, at least, if the receiver’s environment supports a `get_allocator` query. Supporting the receiver’s allocator isn’t always possible: the used allocator type needs to be known when the coroutine is created. At that time the receiver isn’t known, yet.”

This is D4007’s thesis stated as a US national body comment.

US 253 ([cplusplus/nbballot#961](https://github.com/cplusplus/nbballot/issues/961) (<https://github.com/cplusplus/nbballot/issues/961>), LWG4333): “Allow use of arbitrary allocators for coroutine frame.” Status: needs wording.

C.7 Kona Algorithm Customisation Straw Poll

At Kona (November 2025), LEWG reviewed P3826’s proposed fix for algorithm customisation:

Option	Description	F	A
1	Remove all of the C++26 <code>std::execution</code>	3	30
2	Remove all of the customisable sender algorithms for C++26	12	14
3	Remove sender algorithm customisation (early & late)	17	5
4	Fix customisations now (only late CPs)	21	7
5	Ship as-is and fix algorithm customisation in a DR	6	21

Attendance: 44.

Option 4 (fix now) won the most support. The resolution went through three revisions before eventual acceptance in January 2026.

C.8 Outstanding LWG Defects

Two Priority 1 defects remain open as of the Kona meeting:

- **LWG 4206:** `connect_result_t` unconstrained, causing hard errors instead of SFINAE-friendly failures.
- **LWG 4368:** Dangling-reference vulnerability in `transform_sender` (stack-use-after-scope) - returns `xvalue` to a dead temporary, potential undefined behaviour.

Additional open defects:

- **LWG 4190:** `completion-signatures-for` specification is recursive - a circular dependency that cannot be satisfied.
- **LWG 4215:** `run_loop::finish` should be `noexcept` - throwing causes `sync_wait` to hang forever.
- **LWG 4355:** `connect-awaitable()` should mandate receiver completion-signals.
- **LWG 4356:** `connect()` should use `get_allocator(get_env(rcvr))` - directly relevant to the allocator timing issue.

Priority 1 issues in an approved feature are a concrete signal that the specification has not yet stabilised. These are not obscure corner cases but defects in core mechanisms (`connect` and `transform_sender`) that every sender/receiver program exercises.

Appendix D - Post-Approval Modification Catalogue

The following tables list all WG21 papers identified as fixing, reworking, removing, or completing missing functionality in `std::execution` (P2300) after its approval for C++26. Papers that extend the framework into new domains (e.g., networking) are excluded.

Removals

Paper	Title	Author(s)	Date	Status	Change
P3187R1	Remove <code>ensure_started</code> and <code>start_detached</code> from P2300	Lewis Baker, Eric Niebler	2024-10-15	Adopted	Removes two algorithms that dynamically allocate with no allocator customization and break structured concurrency.
P3682R0	Remove <code>std::execution::split</code>	Eric Niebler	2025-02-04	Adopted (Sofia)	Removes <code>split</code> due to incorrect description of its purpose and problematic semantics.

Major Design Reworks

Paper	Title	Author(s)	Date	Status	Change
P2855R1	Member customization points for Senders and Receivers	Ville Voutilainen	2024-03-18	Adopted	Replaces <code>tag_invoke</code> -based ADL customization with member functions - a breaking change.
P2999R3	Sender Algorithm Customization	Eric Niebler	2024-04-16	Adopted	Removes ADL-based customization of sender algorithms in favour of member-function customization on a domain object.
P3303R1	Fixing Lazy Sender Algorithm Customization	Eric Niebler	2024-10-15	Adopted	Fixes gross oversight in P2999 where wording changes that implement the approved design were missing.
P3175R3	Reconsidering the <code>std::execution::on</code> algorithm	Eric Niebler	2024-10-15	Adopted	Renames <code>on</code> to <code>starts_on</code> and <code>transfer</code> to <code>continues_on</code> because usage revealed a gap between users' expectations and actual behaviour. Also fixes a bug in <code>get_scheduler</code> .
P3557R3	High-Quality Sender Diagnostics with Constexpr Exceptions	Eric Niebler	2025-06-10	Adopted (Sofia)	Reworks <code>get_completion_signatures</code> from member function to static constexpr function template. Adds <code>dependent_sender</code> concept.
P3570R2	Optional variants in sender/receiver	Lewis Baker	2025-06-14	Adopted - forwarded to LWG for C++26.	Adds <code>get_await_completion_adapter</code> for coroutine users.
P3718R0	Fixing Lazy Sender Algorithm Customization, Again	Eric Niebler	2025-07-24	In Progress - open, bumped from 2025-telecon	Further fixes to the lazy customization mechanism after P3303.

Paper	Title	Author(s)	Date	Status	Change
P3826R3	Fix Sender Algorithm Customization	Eric Niebler	2025-11-14	In Progress - open, 2026-telecon milestone. Title evolved from “Defer... to C++29” (R0) to “Fix or Remove...” (R1) to “Fix...” (R3). Linked to 5 NB comments. Under active LEWG review.	to 2026-telecon milestone. Third paper attempting to fix lazy customization (after P2999, P3303). Linked to NB comment.
P3927RO	Support for Parallel Bulk Execution <code>task_scheduler</code>	Eric Niebler	2026-01-17	In Progress - open, 2026-telecon milestone. January 2026 mailing. Not yet reviewed in telecon. Implemented in NVIDIA CCCL.	Proposes deferring algorithm customization features that cannot be fixed in time for C+ +26. Fixes <code>task_scheduler</code> not parallelizing bulk work when wrapping a <code>parallel_scheduler</code> .

Wording Fixes and Corrections

Paper	Title	Author(s)	Date	Status	Change
P3396R1	<code>std::execution</code> wording fixes	Eric Niebler	2024-11-22	Adopted	Omnibus paper addressing multiple wording issues: <code>run_loop</code> preconditions, environment ownership, scheduler concept inconsistencies.
P3388R3	When Do You Know <code>connect</code> Doesn't Throw?	Ville Voutilainen	2025-02-14	Adopted	Fixes incorrect <code>noexcept</code> clause of the constructor of <code>basic-state</code> .
P3914R0	Assorted NB comment resolutions for Kona 2025	Various	2025-11-07	In Progress - omnibus NB comment resolution paper. Sections 2.2-2.5 address <code>std::execution</code> . Individual resolutions adopted piecemeal.	Addresses national body comments on the C++26 CD related to <code>std::execution</code> .
P3887R1	Make <code>when_all</code> a Ronseal Algorithm	Robert Leahy	2025-11-07	Adopted - forwarded at Kona (SF:10/F:5/N:0/A:0/SA:0), wording merged into draft (Dec 2025).	Fixes <code>when_all</code> stop-request handling - removes unnecessary stop-detection complexity that made the algorithm's behaviour surprising.
P3940R0	Rename concept tags for		2025-12-15	In Progress - open, 2026-telecon	Renames concept tag types

Paper	Title	Author(s)	Date	Status	Change
C++26:	<code>sender_t</code> to <code>sender_tag</code>	Arthur O'Dwyer, Yi'an Ye		milestone. Post-Kona mailing. Not yet reviewed in telecon.	(<code>sender_t</code> to <code>sender_tag</code> , etc.) for naming consistency - another post-approval naming correction.

Missing Functionality

Paper	Title	Author(s)	Date	Status	Change
P3425R1	Reducing operation-state sizes for subobject child operations	Eric Niebler	2024-11-19	LEWG Wroclaw (Nov 2024) strong consensus.	Design approved - Optimisation saving 8 bytes per nesting level - performance fix for deeply nested sender expressions. C++26-targeted.
P3284R4	<code>write_env</code> and <code>unstoppable</code> Sender Adaptors	Eric Niebler	2025-02-14	Adopted (Sofia)	Adds missing sender adaptors for modifying execution environments.
P3685R0	Rename <code>async_scope_token</code>	Ian Petersen, Jessica Wong	2025-04-09	Adopted	Renames <code>async_scope_token</code> to <code>scope_token</code> for clarity.
P3706R0	Rename <code>join</code> and <code>nest</code> in async scope proposal	Ian Petersen, Jessica Wong	2025-04-09	Adopted	Renames <code>nest</code> to <code>associate</code> because original names were misleading.
P3325R5	A Utility for Creating Execution Environments	Eric Niebler	2025-05-22	Adopted	Adds <code>prop</code> and <code>env</code> class templates for creating and manipulating

Paper	Title	Author(s)	Date	Status	Change
P2079R10	Parallel scheduler	Lee Howes	2025-06-02	Adopted (Sofia)	environments - fundamental infrastructure that was absent.
P3149R11	<code>async_scope</code>	Ian Petersen, Jessica Wong, Kirk Shoop, et al.	2025-06-02	Adopted (Sofia)	Provides <code>system_context</code> and <code>system_scheduler</code> - a basic execution context needed to actually run code.
P3164R4	Early Diagnostics for Sender Expressions	Eric Niebler	2025-06-02	Adopted	Provides the async scope abstraction needed for safe non-sequential concurrency - replacing the removed <code>ensure_started</code> / <code>start_detached</code> .
P3552R3	Add a Coroutine Task Type	Dietmar Kuhl, Maikel Nadolski	2025-06-20	Adopted (Sofia)	Moves diagnosis of invalid sender expressions to construction time rather than connection time.
P3815R1	Add <code>scope_association</code> concept to P3149	Jessica Wong, Ian Petersen	2025-09-12	Adopted - closed Dec 2025. NB comment resolution.	Adds <code>std::execution::task</code> - the coroutine type that users need to use the framework. Adopted at Sofia with 29 abstentions and 11 against (77-11-29).
					Adds missing <code>scope_association</code> concept needed by the async scope facility.

Post-Adoption Issues

Paper	Title	Author(s)	Date	Status	Change
P3433R1	Allocator Support for Operation States	Dietmar Kuhl	2024-10-17	Plenary-approved - closed. LEWG Wroclaw (Nov 2024) approved design with strong consensus (SF:4/ F:6/N:0/A:0/ SA:0). Wording merged into draft.	Identifies allocator gaps in <code>ensure_started</code> , <code>split</code> , and <code>start_detached</code> operation states - confirming allocator support was missing from the original design.
P3481R5	<code>std::execution::bulk()</code> issues	Lucian Radu, Teodorescu, Lewis Baker, Ruslan Arutyunyan	2024-10-17	Plenary-approved - closed. Five revisions. SG1 Wroclaw (Nov 2024) achieved unanimous consent on splitting bulk into <code>bulk</code> , <code>bulk_chunked</code> , and <code>bulk_unchunked</code> . Wording merged into draft.	Addresses outstanding issues with the bulk algorithm; required five revisions and addition of two new API variants (<code>bulk_chunked</code> , <code>bulk_unchunked</code>).
P3796R1	Coroutine Task Issues	Dietmar Kuhl	2025-07-24	In Progress - open, 2026-telecon milestone. Under active LEWG telecon review (Aug-Sep 2025); some sections achieved consensus,	Collects issues discovered after the task type was forwarded, including <code>Unhandled_stopped</code> missing <code>noexcept</code> , wording issues, and performance concerns.

Paper	Title	Author(s)	Date	Status	Change
P3801R0	Concerns about the design of <code>std::execution::task</code>	Jonathan Wakely	2025-07-24	<p>others still pending. Multiple LWG issues opened (4329-4332, 4344). Linked to NB comments.</p> <p>In Progress - open, 2026-telecon milestone. LEWG telecon review (2025-08-26) reached “no consensus” on the core stack overflow issue; “consensus against” treating dangling reference concern as C+ +26 blocker. Linked to NB comment.</p>	<p>Documents significant concerns including stack overflow risk due to lack of symmetric transfer support.</p>

LWG Issues

Issue	Title	Date	Status	Change
LWG 4190	<code>completion-signatures-for</code> specification is recursive	2025-01-02	Open - circular dependency in spec.	Specification defect: recursive definition makes the requirement impossible to implement.
LWG 4206	<code>connect_result_t</code> should be constrained with <code>sender_to</code>	2025-02-04	Open - Priority 1.	Unconstrained alias causes hard errors instead of SFINAE-friendly failures.

2025-02-13

Issue	Title	Date	Status	Change
LWG 4215	<code>run_loop::finish</code> should be <code>noexcept</code>		Open - correctness bug.	Throwing <code>finish()</code> causes <code>sync_wait</code> to hang forever.
LWG 4260	Query objects must be default constructible	2025-05-07	Resolved (Kona 2025)	CPO constructors were not mandated <code>noexcept</code> .
LWG 4355	<code>connect-awaitable()</code> should mandate receiver completion-signals	2025-08-27	Open	Redundant <code>requires</code> -clause should defer to parent <code>Mandates</code> .
LWG 4356	<code>connect()</code> should use <code>get_allocator(get_env(rcvr))</code>	2025-08-27	Open - directly relevant to allocator timing issue.	<code>connect-awaitable</code> should respect the receiver's allocator.
LWG 4358	<code>[exec.as.awaitable]</code> uses Preconditions when should be constraint	2025-08-27	Resolved (Kona 2025)	Incorrectly-written preconditions should be <code>Mandates</code> .
LWG 4360	<code>awaitable-sender</code> concept should qualify <code>awaitable-receiver</code>	2025-08-27	Resolved (Kona 2025)	Ambiguous type reference in the <code>awaitable-sender</code> concept.
LWG 4368	Potential dangling reference from <code>transform_sender</code>	2025-08-31	Open - Priority 1. Stack-use-after-scope vulnerability.	Returns xvalue to a dead temporary - potential undefined behaviour.
LWG 4369	<code>check-types</code> for <code>upon_error</code> and <code>upon_stopped</code> is wrong	2025-08-31	Resolved (Kona 2025)	Uses <code>set_value_t</code> where <code>set_error_t</code> / <code>set_stopped_t</code> should be used.
LWG 4336	<code>bulk</code> vs. <code>task_scheduler</code>	2025-10-23	Open - NB comment. C++26-targeted.	<code>task_scheduler</code> does not parallelise bulk work - dispatching <code>bulk</code> via a <code>task_scheduler</code> wrapping <code>parallel_scheduler</code> serialises execution. P3927R0 proposes the fix.

Allocator-Related NB Comments (Kona 2025)

NB Comment	Title	Status
US 255 ((https://github.com/cplusplus/nbballot/issues/959) (LWG4335))	Use allocator from receiver's environment	Needs wording
US 253 ((https://github.com/cplusplus/nbballot/issues/961) (LWG4333))	Allow use of arbitrary allocators for coroutine frame	Needs wording

Total: 30 papers, 11 LWG issues, and 2 NB comments - 43 items modifying a single feature after its approval.

References

- P2079R10 (<https://wg21.link/p2079r10>) Lee Howes. “Parallel scheduler.” 2025-06-02.
- P2300R4 (<https://wg21.link/p2300r4>) Michal Dominiak, et al. “std::execution.” 2022-01-18.
- P2300R10 (<https://wg21.link/p2300r10>) Michal Dominiak, et al. “std::execution.” 2024-07-16.
- P2762R2 (<https://wg21.link/p2762r2>) Dietmar Kuhl. “Sender/Receiver Interface For Networking.” 2023-10-15.
- P2855R1 (<https://wg21.link/p2855r1>) Ville Voutilainen. “Member customization points for Senders and Receivers.” 2024-03-18.
- P2999R3 (<https://wg21.link/p2999r3>) Eric Niebler. “Sender Algorithm Customization.” 2024-04-16.
- P3149R11 (<https://wg21.link/p3149r11>) Ian Petersen, Jessica Wong, Kirk Shoop, et al. “async_scope.” 2025-06-02.
- P3175R3 (<https://wg21.link/p3175r3>) Eric Niebler. “Reconsidering the std::execution::on algorithm.” 2024-10-15.
- P3187R1 (<https://wg21.link/p3187r1>) Lewis Baker, Eric Niebler. “Remove ensure_started and start_detached from P2300.” 2024-10-15.
- P3433R1 (<https://wg21.link/p3433r1>) Dietmar Kuhl. “Allocator Support for Operation States.” 2025-06-18.
- P3552R3 (<https://wg21.link/p3552r3>) Dietmar Kuhl, Maikel Nadolski. “Add a Coroutine Task Type.” 2025-06-20.
- P3557R3 (<https://wg21.link/p3557r3>) Eric Niebler. “High-Quality Sender Diagnostics with Constexpr Exceptions.” 2025-06-10.
- P3796R1 (<https://wg21.link/p3796r1>) Dietmar Kuhl. “Coroutine Task Issues.” 2025-07-24.
- P3801R0 (<https://wg21.link/p3801r0>) Jonathan Wakely. “Concerns about the design of std::execution::task.” 2025-07-24.
- P3826R3 (<https://wg21.link/p3826r3>) Eric Niebler. “Fix Sender Algorithm Customization.” 2026-01-05.

- [P3927R0](https://wg21.link/p3927r0) (https://wg21.link/p3927r0) Eric Niebler. “task_scheduler Support for Parallel Bulk Execution.”
2026-01-17.
- [P4003R0](https://wg21.link/p4003r0) (https://wg21.link/p4003r0) Vinnie Falco. “IoAwaitables: A Coroutines-First Execution Model.”
2026-01-21.